Learning to Search for Targets
## with Deep Reinforcement Learning

Oskar Lundin

Linköping University

June 9, 2022

Outline

Description

Learned autonomous search for a set of targets in a visual environment
with a camera.

- ▶ Camera perceives limited region of environment.
- ▶ Moving camera changes visible region.
- ▶ Detect when targets are visible.
- ▶ Locate targets in minimum time.
- ▶ Learn control from sample scenarios.
- ▶ Use deep reinforcement learning.
- ▶ Focus on search paths, simple detection.

Motivation

- ▶ Applications in search and rescue, surveillance, home assistance, etc.
- ▶ Autonomous systems may reduce risk and cost.
- ▶ Advantages of learning systems over handcrafted ones:
  - ▶ May find better solutions (deep RL: Atari [1], Go [2], StarCraft II [3]).
  - ▶ Applicable as long as data is available.
 - ▶ Sufficient to describe problem.

Aim

- ▶ Utilize structure in environments:
  - ▶ Books are in bookshelves, cars on roads...
  - ▶ Targets can be spread out/close together...
- ▶ Learn distribution of targets from training samples.
  - ▶ Realistically limited training samples available.
  - ▶ Generalize to similar unseen search tasks.
- ▶ Remember features of explored environment to:
  - ▶ Avoid searching regions twice.
  - ▶ Prioritize promising regions.

Research Questions

1. How can an agent that learns to intelligently search for targets be implemented with deep reinforcement learning?
2. How does the learning agent compare to random, greedy, exhaustive and human searchers with prior knowledge of the searched scenes?
3. How well does the learning agent generalize from a limited number of training samples to unseen in-distribution search scenarios?

Reinforcement Learning I

- ▶ Learn from interaction how to achieve a goal.
- ▶ Partially Observable Markov Decision Process [4]:
  - ▶ *Agent* interacts with *environment* over discrete time steps $t = 0, 1, 2 \ldots, T$.
  - ▶ New state s_{t+1} depends on history $a_0, o_1, r_1, \ldots, a_{t-1}, o_t, r_t$.
  - ▶ Agent usually maintains internal state $\rightarrow$ memory.

Reinforcement Learning II

- ▶ Policy $\pi(a|s)$ is a mapping from states to action probabilities.
- ▶ Find policy that maximizes expected future reward
  $\mathbb{E}\left[\sum_{k=0}^{T} \gamma^{k-t-1} r_k\right]$.
- ▶ Reward signal is often a design parameter.
- ▶ Deep reinforcement learning: approximate π with deep neural networks.

Related Work

- ▶ Visual attention:
  - ▶ Sequential focus points for foveated vision [5].
- ▶ Visual navigation:
  - ▶ Solve random mazes [6].
  - ▶ Find target object in indoor scenes [7].
- ▶ Object detection:
  - ▶ Region proposals for object localization [8].
  - ▶ Contextual reasoning over spatial layout in scenes [9].
  - ▶ Anatomical landmark detection in medical images [10].

Problem Statement

- ▶ Agent searches scene $S \subset \mathbb{R}^d$.
- ▶ Agent perceives view $V \subset S$.
- ▶ View can be transformed to new subspace.
- ▶ Targets in scene $\{t_0, \ldots t_n\}$, $t_i \in S$.
- ▶ Indicate when targets are visible, i.e. $V \cup T \neq \varnothing$.
- ▶ Goal:
  - ▶ Maximize probability of finding all targets.
  - ▶ Minimize cost in time.
 - ▶ NP-complete [11]).

Environments

- ▶ Three simulated environments.
- ▶ Find three targets in less than 1 000 steps.
- ▶ Target probability correlated with scene appearance.
- ▶ Procedurally generated, conditioned on seed.
- ▶ New seed after each finished search.

Observation, Action and Reward

- Observations $o_t = \langle x_t, p_t \rangle$, where
  - $x_t \in \mathbb{R}^{3 \times 64 \times 64}$ is an RGB image,
  - $p_t \in \{0, \ldots, H\} \times \{0, \ldots, W\}$ is the camera position.
- Actions $a_t \in \{\texttt{INDICATE}, \texttt{UP}, \texttt{DOWN}, \texttt{LEFT}, \texttt{RIGHT}\}$, where
  - $\texttt{INDICATE}$ identifies targets, and
  - $\texttt{UP}, \texttt{DOWN}, \texttt{LEFT}, \texttt{RIGHT}$ move camera.
- Reward $r_t = h - 0.01 + 0.005d + 0.005e$ where
  - $h = |T \cap V|$ if $a_t = \texttt{INDICATE}$, else 0.
  - $d = 1$ if a_t moves closer to nearest target, else 0.
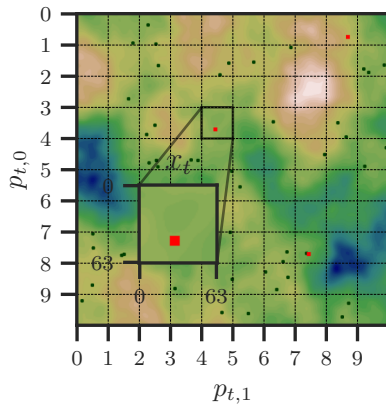 - $e = 1$ if a_t moves to new position, else 0.

Environment I: Gaussian

▶ Three gaussian kernels with random center.

▶ Sum of kernels = blue color intensity, probability of targets.
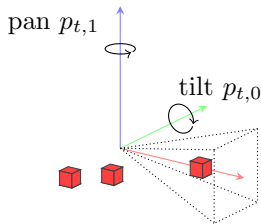
▶ Agent should prioritize blue regions.
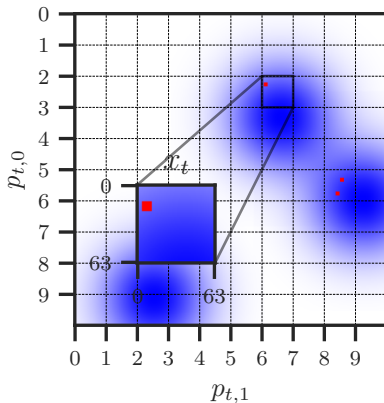
Environment II: Terrain

▶ Terrain seen from above (e.g. UAV).

▶ Targets between ocean and mountains.

▶ More realistic, higher variance.

Environment III: Camera
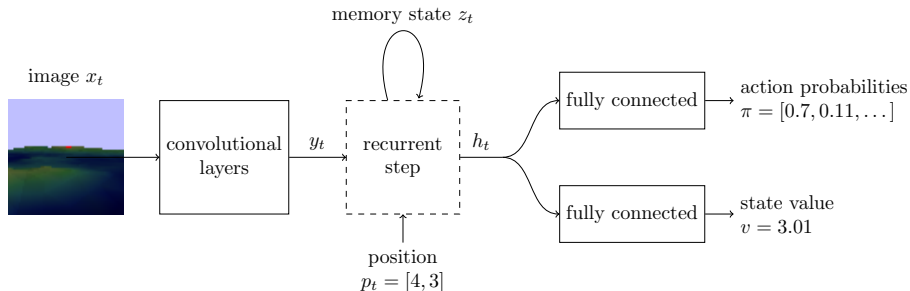
▶ Terrain seen from perspective projection camera.

▶ Variance in target appearance.

▶ Moving actions control pan and tilt.

    ▶ 20 pan angle steps.

    ▶ 10 tilt angle steps.

x_t

Approach

- ▶ Function approximation with deep neural networks.
    - ▶ Policy $\pi(a|s, \theta)$.
    - ▶ Value $v_\pi(s, \theta)$ (predicts future reward).
- ▶ Training procedure:
    1. Collect interactions with environment.
    2. Compute loss $\mathcal{L}(\theta)$.
    3. Optimize $\mathcal{L}$ wrt θ.
    4. Repeat. . .
- ▶ Use proximal policy optimization [12].
    - ▶ Clipping loss function $\mathcal{L}_{\text{clip}}$.
    - ▶ RL algorithm from 2017.
    - ▶ Stable performance, relatively little tuning [13].

Architecture

memory state z_t

image x_t

convolutional layers $\xrightarrow{y_t}$ recurrent step $\xrightarrow{h_t}$

fully connected → action probabilities $\pi = [0.7, 0.11, \dots]$

fully connected → state value $v = 3.01$

position $p_t = [4, 3]$

Memory

- *Agent should remember visual features and associate them with their spatial location.*
- Two memory variants:
    1. Temporal memory (long short-term memory [14]):
        - Previously applied to tasks where memory is required [15, 16, 6, 17].
        - May struggle with remembering over many time steps.
        - Important for exhaustive search and scene understanding.
    2. Spatial memory (inspired by [18]):
        - Matrix with one slot per camera position.
        - Indexed with current position.
        - Stores image representation at each slot.
        - Read whole memory with convolutional layers.

Training

- ▶ Train for 25M time steps.
- ▶ Results reported across 3 training runs.
- ▶ Separate training and test sets.
- ▶ Same hyperparameters in all runs.

Implementation

- ▶ OpenAI Gym environment interface.
- ▶ Custom PPO implementation.
- ▶ PyTorch for models and automatic differentiation.
- ▶ Intel Core i9-10900X CPU.
- ▶ NVIDIA GeForce RTX 2080 Ti GPU.

Experiment I: Search Performance

▶ Compare to simple reference behaviors (baselines).

▶ Test on held out environment samples.

▶ Metrics:
   1. Average search path length.
   2. Average success rate.
   3. Success weighted by inverse path length (SPL) [19].

Definition

$$\text{SPL} = \frac{1}{N} \sum_{i=1}^{N} S_i \frac{l_i}{\max(p_i, l_i)}$$

where N is number of test samples, S_i is binary success indicator, p_i is the taken path length l_i is the shortest path length.

Baselines

Random: randomly samples actions.

Greedy: greedily selects exploring actions (random if none).

Exhaustive: exhaustively covers search space with minimal revisits.

Human: human searcher with knowledge of environment.

Handcrafted: prioritize actions that lead to higher blue intensity (gaussian environment only).

Gaussian Environment

| Agent	SPL	Success	Length
random	0.06 ± 0.01	0.92 ± 0.06	369.07 ± 24.93
greedy	0.17 ± 0.00	1.00 ± 0.00	147.12 ± 2.38
exhaustive	0.21 ± 0.00	1.00 ± 0.00	83.37 ± 2.88
handcrafted	0.33 ± 0.00	1.00 ± 0.00	65.20 ± 1.41
human	0.23 ± 0.03	1.00 ± 0.00	80.97 ± 13.49
temporal	0.24 ± 0.03	0.99 ± 0.01	101.25 ± 13.32
spatial	0.29 ± 0.02	0.99 ± 0.01	72.16 ± 5.97

Environment sample

Random baseline

Environment sample


Greedy baseline

Environment sample


Exhaustive baseline

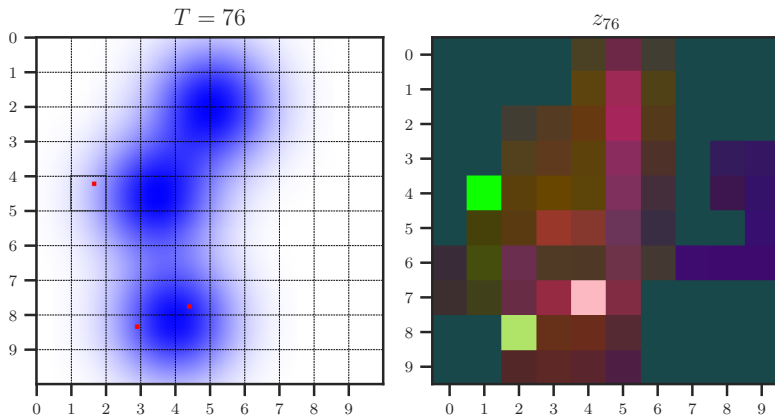Environment sample

Handcrafted baseline

Environment sample

Temporal memory

Environment sample

Spatial memory

PCA decomposition of spatial memory after episode.

Terrain Environment

Agent	SPL	Success	Length
random	0.06 ± 0.01	0.89 ± 0.04	366.05 ± 26.96
greedy	0.17 ± 0.01	1.00 ± 0.00	141.01 ± 2.31
exhaustive	0.22 ± 0.00	1.00 ± 0.00	84.11 ± 0.84
human	0.26 ± 0.02	1.00 ± 0.00	76.73 ± 5.33
temporal	0.25 ± 0.02	1.00 ± 0.01	103.76 ± 11.69
spatial	0.27 ± 0.01	1.00 ± 0.00	79.60 ± 6.88

video 1, video 2, video 3 (spatial)

Camera Environment

Agent	SPL	Success	Length
random	0.04 ± 0.00	0.62 ± 0.03	545.09 ± 56.25
greedy	0.12 ± 0.01	0.97 ± 0.01	255.60 ± 10.44
exhaustive	0.37 ± 0.00	1.00 ± 0.00	67.03 ± 0.00
human	0.68 ± 0.08	1.00 ± 0.00	38.10 ± 5.72
temporal	0.70 ± 0.02	1.00 ± 0.00	42.36 ± 2.05
spatial	0.66 ± 0.03	1.00 ± 0.00	42.90 ± 1.73

video 1, video 2, video 3 (temporal)

Experiment II: Scaling to Larger Search Spaces

▶ Larger search spaces are more difficult.
▶ Stronger demands on memory:
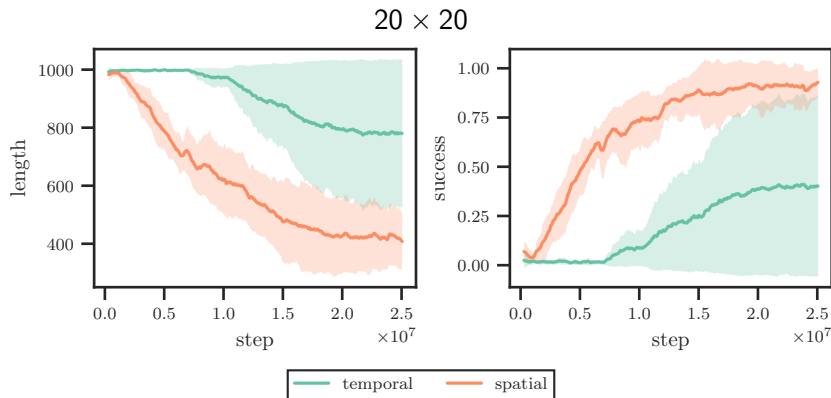 ▶ Remember visited positions.
  ▶ Remember appearance of environment.
▶ Compare memories on 10×10, 15×15, and 20×20 versions of gaussian environment.
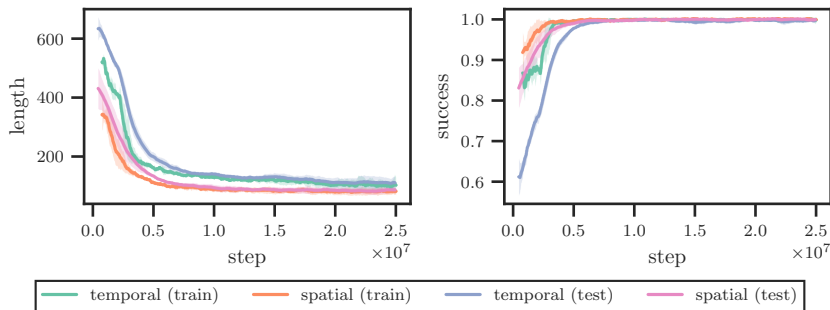
10×10

15×15

20×20
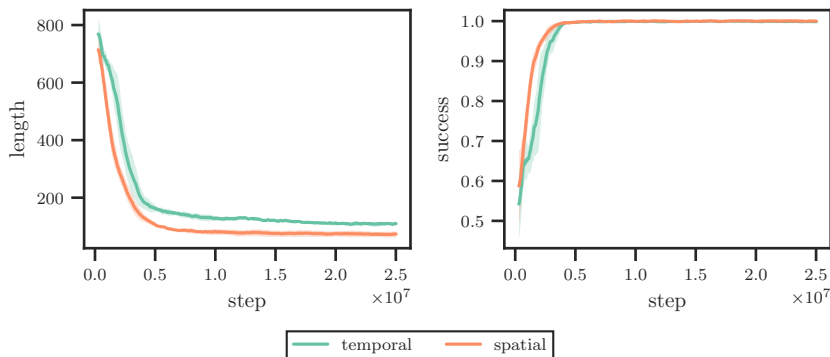
Experiment III: Generalization From Limited Samples

- ▶ Real-world use-cases have limited training samples.
- ▶ Train on 500, 1 000, 5 000 and 10 000 samples of terrain environment.
- ▶ Test on held out samples from full distribution.

10000 samples

5000 samples
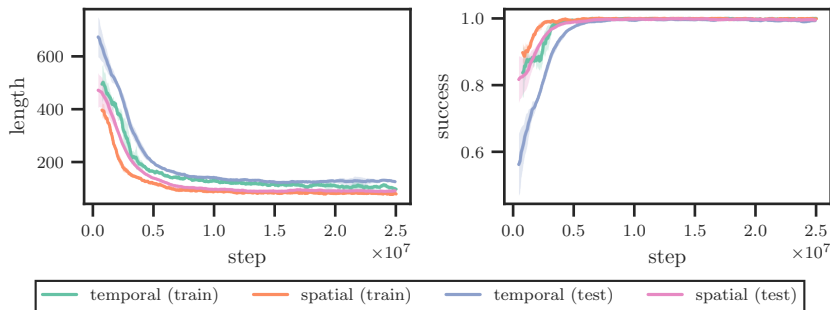
1000 samples
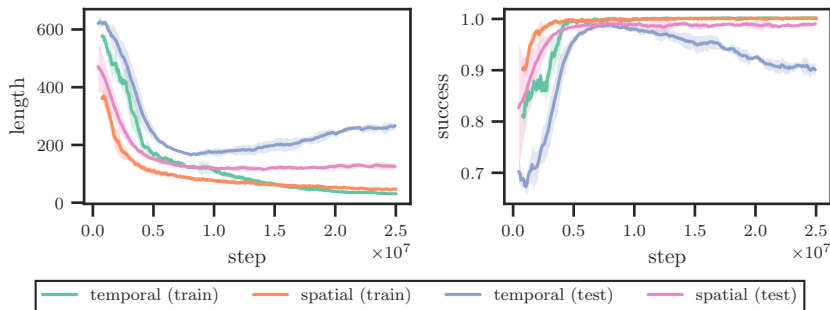
500 samples
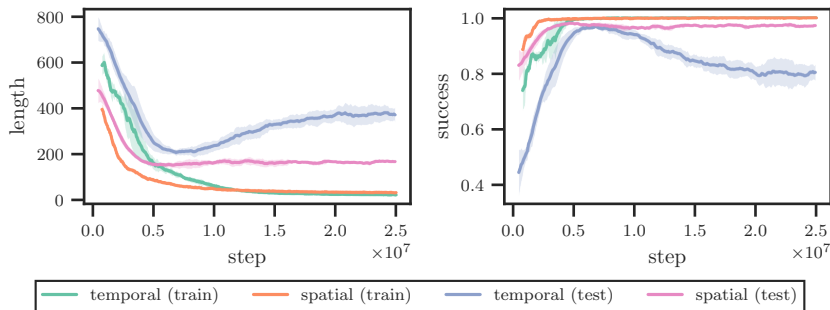
Conclusion

- ▶ General method for visual search with reinforcement learning.
- ▶ Three environments for evaluating visual search agents.
- ▶ Two different neural network architectures.
- ▶ Architecture affects performance, scaling and generalization.
- ▶ One approach comparable to human performance.

Future Work

- ▶ Real-world scenarios.
- ▶ Formal verification (for security-critical applications).
- ▶ Hyperparameter tuning (expensive!).
- ▶ Other RL algorithms.

References I

[1] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves,
     M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik,
     I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, "Human-level
     control through deep reinforcement learning," vol. 518, no. 7540, pp. 529–533.
     Number: 7540 Publisher: Nature Publishing Group.

[2] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. van den Driessche,
     J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, S. Dieleman, D. Grewe,
     J. Nham, N. Kalchbrenner, I. Sutskever, T. Lillicrap, M. Leach, K. Kavukcuoglu,
     T. Graepel, and D. Hassabis, "Mastering the game of go with deep neural networks and
     tree search," vol. 529, no. 7587, pp. 484–489.
     Number: 7587 Publisher: Nature Publishing Group.

[3] O. Vinyals, I. Babuschkin, W. M. Czarnecki, M. Mathieu, A. Dudzik, J. Chung, D. H. Choi,
     R. Powell, T. Ewalds, P. Georgiev, J. Oh, D. Horgan, M. Kroiss, I. Danihelka, A. Huang,
     L. Sifre, T. Cai, J. P. Agapiou, M. Jaderberg, A. S. Vezhnevets, R. Leblond, T. Pohlen,
     V. Dalibard, D. Budden, Y. Sulsky, J. Molloy, T. L. Paine, C. Gulcehre, Z. Wang, T. Pfaff,
     Y. Wu, R. Ring, D. Yogatama, D. Wünsch, K. McKinney, O. Smith, T. Schaul,
     T. Lillicrap, K. Kavukcuoglu, D. Hassabis, C. Apps, and D. Silver, "Grandmaster level in
     StarCraft II using multi-agent reinforcement learning," vol. 575, no. 7782, pp. 350–354.
     Number: 7782 Publisher: Nature Publishing Group.

References II

[4] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, "Planning and acting in partially observable stochastic domains," vol. 101, no. 1, pp. 99–134.

[5] V. Mnih, N. Heess, A. Graves, and k. kavukcuoglu, "Recurrent models of visual attention," in *Advances in Neural Information Processing Systems*, vol. 27, Curran Associates, Inc.

[6] P. Mirowski, R. Pascanu, F. Viola, H. Soyer, A. J. Ballard, A. Banino, M. Denil, R. Goroshin, L. Sifre, K. Kavukcuoglu, D. Kumaran, and R. Hadsell, "Learning to navigate in complex environments,"

[7] Y. Zhu, R. Mottaghi, E. Kolve, J. J. Lim, A. Gupta, L. Fei-Fei, and A. Farhadi, "Target-driven visual navigation in indoor scenes using deep reinforcement learning," in *2017 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 3357–3364. 474 citations (Crossref) [2022-05-19].

[8] J. C. Caicedo and S. Lazebnik, "Active object localization with deep reinforcement learning,"

[9] X. Chen and A. Gupta, "Spatial memory for context reasoning in object detection," in *2017 IEEE International Conference on Computer Vision (ICCV)*, pp. 4106–4116, IEEE.

References III

[10] F.-C. Ghesu, B. Georgescu, Y. Zheng, S. Grbic, A. Maier, J. Hornegger, and D. Comaniciu, "Multi-scale deep reinforcement learning for real-time 3d-landmark detection in CT scans," vol. 41, no. 1, pp. 176–189.

[11] A. Andreopoulos and J. K. Tsotsos, "A theory of active object localization," in *2009 IEEE 12th International Conference on Computer Vision*, pp. 903–910.
15 citations (Crossref) [2022-05-19] ISSN: 2380-7504.

[12] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms,"

[13] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, "Deep reinforcement learning that matters," vol. 32, no. 1.
Number: 1.

[14] S. Hochreiter and J. Schmidhuber, "Long short-term memory," vol. 9, no. 8, pp. 1735–1780.
Conference Name: Neural Computation.

[15] M. Hausknecht and P. Stone, "Deep recurrent q-learning for partially observable MDPs,"

[16] V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. P. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, "Asynchronous methods for deep reinforcement learning,"

References IV

[17] S. Gupta, V. Tolani, J. Davidson, S. Levine, R. Sukthankar, and J. Malik, "Cognitive mapping and planning for visual navigation,"

[18] E. Parisotto and R. Salakhutdinov, "Neural map: Structured memory for deep reinforcement learning,"

[19] P. Anderson, A. Chang, D. S. Chaplot, A. Dosovitskiy, S. Gupta, V. Koltun, J. Kosecka, J. Malik, R. Mottaghi, M. Savva, and A. R. Zamir, "On evaluation of embodied navigation agents,"