


Learning to Search for Targets

with Deep Reinforcement Learning

Oskar Lundin

Linköping University

June 12, 2022



Outline

Introduction

- Motivation

- Aim

- Research Questions

Theory

- Reinforcement Learning

- Related Work

Method

- Environments

- Approach

Experiments

- Search Performance

- Scaling to Larger Search Spaces

- Generalization From Limited Samples

Conclusion

- Future Work


Description

Learned autonomous search for a set of targets in a visual environment with a camera.

- ▶ Camera perceives limited region of environment.
- ▶ Moving camera changes visible region.
- ▶ Detect when targets are visible.
- ▶ Locate targets in minimum time.
- ▶ Learn control from sample scenarios.
- ▶ Use deep reinforcement learning.
- ▶ Focus on search behavior, simple detection.

Problem Statement

- ▶ Searched scene $S \subset \mathbb{R}^d$.
- ▶ Perceived view $V \subset S$ in the form of an image.
- ▶ View can be transformed to new subspace at a cost.
- ▶ Targets in scene $\{t_0, \dots, t_n\}$, $t_i \in S$.
- ▶ Detect when targets are visible, i.e. $V \cap T \neq \emptyset$.
- ▶ Goal:
 - ▶ Maximize probability of finding all targets.
 - ▶ Minimize cost (time).
 - ▶ NP-complete [1].



Motivation

- ▶ Applications in search and rescue, surveillance, home assistance, etc.
- ▶ Autonomous systems may reduce cost and time.
- ▶ Learning vs. handcrafted systems:
 - ▶ May find better solutions (deep RL: Atari [2], Go [3], StarCraft II [4]).
 - ▶ Applicable as long as data is available.
 - ▶ Guarantees and understandability.

Aim


- ▶ Utilize structure in environments:
 - ▶ Books are in bookshelves, cars on roads...
 - ▶ Targets can be spread out/close together...
- ▶ Learn distribution of targets from training samples.
 - ▶ Realistically limited training samples available.
 - ▶ Generalize to similar unseen search scenarios.
- ▶ Remember features of explored environment to:
 - ▶ Avoid searching regions twice.
 - ▶ Prioritize promising regions.

Research Questions

1. How can an agent that learns to intelligently search for targets be implemented with deep reinforcement learning?
2. How does the learning agent compare to random, greedy, exhaustive and human searchers?
3. How well does the learning agent generalize from a limited number of training samples to unseen in-distribution search scenarios?

Reinforcement Learning I

- ▶ Learn from interaction how to achieve a goal.
- ▶ Partially Observable Markov Decision Process [5]:
 - ▶ *Agent* interacts with *environment* over discrete time steps $t = 0, 1, 2 \dots, T$.
 - ▶ New state s_{t+1} depends on history $a_0, o_1, r_1, \dots, a_{t-1}, o_t, r_t$.
 - ▶ Agent usually maintains internal state \rightarrow memory.



Reinforcement Learning II

- ▶ Policy $\pi(a|s)$ defines agent's behavior.
- ▶ Find policy that maximizes expected future reward $\mathbb{E} \left[\sum_{k=0}^T \gamma^{k-t-1} r_k \right]$.
- ▶ There are several different algorithms.
- ▶ Reward signal is often a design parameter.
- ▶ Deep reinforcement learning: approximate π with deep neural networks.

Related Work

- ▶ Visual attention:
 - ▶ Sequential focus points for foveated vision [6].
- ▶ Visual navigation:
 - ▶ Solve random mazes [7].
 - ▶ Find target object in indoor scenes [8].
- ▶ Object detection:
 - ▶ Region proposals for object localization [9].
 - ▶ Anatomical landmark detection in medical images [10].

Environments


- ▶ Three simulated environments.
- ▶ Find three targets in less than 1 000 steps.
- ▶ There is some structure that can be utilized to find targets quicker.
- ▶ Procedurally generated, conditioned on seed.
- ▶ New seed after each finished search.

Observation, Action and Reward

- ▶ Observations $o_t = \langle x_t, p_t \rangle$, where
 - ▶ $x_t \in \mathbb{R}^{3 \times 64 \times 64}$ is an RGB image,
 - ▶ $p_t \in \{0, \dots, H\} \times \{0, \dots, W\}$ is the camera position.
- ▶ Actions $a_t \in \{\text{INDICATE}, \text{UP}, \text{DOWN}, \text{LEFT}, \text{RIGHT}\}$, where
 - ▶ INDICATE identifies targets, and
 - ▶ UP, DOWN, LEFT, RIGHT move camera.
- ▶ Reward $r_t = h - 0.01 + 0.005d + 0.005e$ where
 - ▶ $h = |T \cap V|$ if $a_t = \text{INDICATE}$, else 0.
 - ▶ $d = 1$ if a_t moves closer to nearest target, else 0.
 - ▶ $e = 1$ if a_t moves to new position, else 0.


Environment I: Gaussian

- ▶ Three gaussian kernels with random center.
- ▶ Sum of three gaussian kernels = blue color intensity.
- ▶ More blue \rightarrow higher target probability.
- ▶ Agent should prioritize blue regions.





Environment II: Terrain

- ▶ Terrain seen from above (e.g. UAV).
- ▶ Targets between ocean and mountains.
- ▶ More realistic, higher variance.



Environment III: Camera

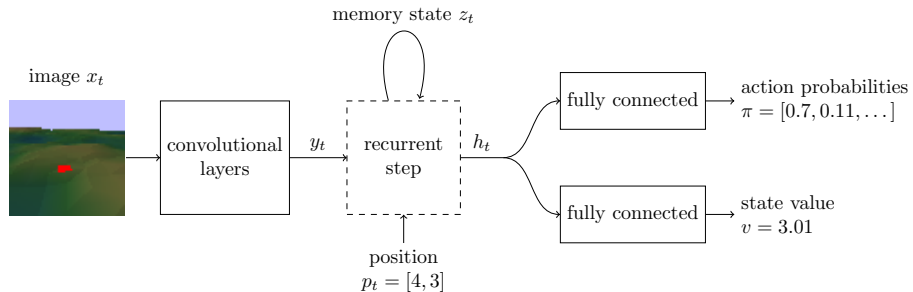
- ▶ Terrain seen from perspective projection camera.
- ▶ Variance in target appearance.
- ▶ Moving actions control pan and tilt.
 - ▶ 20 pan angle steps.
 - ▶ 10 tilt angle steps.



Approach

- ▶ Function approximation with deep neural networks.
 - ▶ Policy $\pi(a|s, \theta)$.
 - ▶ Value $v_\pi(s, \theta)$ (predicts future reward).
- ▶ Training procedure:
 1. Collect interactions with environment.
 2. Compute loss $\mathcal{L}(\theta)$.
 3. Optimize \mathcal{L} wrt θ .
 4. Repeat until π is good.
- ▶ Use proximal policy optimization [11].
 - ▶ RL algorithm from 2017.
 - ▶ Stable performance, relatively little tuning [12].

Architecture



Memory

- ▶ *Agent should remember visual features and associate them with their spatial location.*
- ▶ Two memory variants:
 1. Temporal memory (long short-term memory [13]):
 - ▶ Previously applied successfully to tasks where memory is required [14, 15].
 - ▶ How long sequences can be remembered?
 2. Spatial memory (inspired by [16] and [17]):
 - ▶ Feature map with one slot per camera position.
 - ▶ Indexed with current position.
 - ▶ Stores image representation at each slot.
 - ▶ Read whole memory with convolutional layers.

Training

- ▶ Train for 25M time steps.
- ▶ Results reported across 3 training runs.
- ▶ Separate training and test sets.

Implementation

- ▶ OpenAI Gym environment interface.
- ▶ Custom proximal policy optimization implementation.
- ▶ PyTorch for models and automatic differentiation.
- ▶ Intel Core i9-10900X CPU.
- ▶ NVIDIA GeForce RTX 2080 Ti GPU.

Experiment I: Search Performance

- ▶ Compare to simple reference behaviors (baselines).
- ▶ Fixed test set from each environment.
- ▶ Metrics:

1. Average search path length.
2. Average success rate.
3. Success weighted by inverse path length (SPL) [18].

With N test samples, S_i as a binary success indicator, p_i as the taken search path length l_i is the shortest search path length:

$$\text{SPL} = \frac{1}{N} \sum_{i=1}^N S_i \frac{l_i}{\max(p_i, l_i)}$$

Baselines

Random: randomly samples actions.

Greedy: greedily selects exploring actions (random if none).

Exhaustive: exhaustively covers search space with minimal revisits.

Human: human searcher with knowledge of environment.

Handcrafted: prioritize actions that lead to higher blue intensity
(gaussian environment only).

Gaussian Environment

Agent	SPL	Success	Length
random	0.06 ± 0.01	0.92 ± 0.06	369.07 ± 24.93
greedy	0.17 ± 0.00	1.00 ± 0.00	147.12 ± 2.38
exhaustive	0.21 ± 0.00	1.00 ± 0.00	83.37 ± 2.88
handcrafted	0.33 ± 0.00	1.00 ± 0.00	65.20 ± 1.41
human	0.23 ± 0.03	1.00 ± 0.00	80.97 ± 13.49
temporal	0.24 ± 0.03	0.99 ± 0.01	101.25 ± 13.32
spatial	0.29 ± 0.02	0.99 ± 0.01	72.16 ± 5.97

video 1, video 2, video 3.

Terrain Environment

Agent	SPL	Success	Length
random	0.06 ± 0.01	0.89 ± 0.04	366.05 ± 26.96
greedy	0.17 ± 0.01	1.00 ± 0.00	141.01 ± 2.31
exhaustive	0.22 ± 0.00	1.00 ± 0.00	84.11 ± 0.84
human	0.26 ± 0.02	1.00 ± 0.00	76.73 ± 5.33
temporal	0.25 ± 0.02	1.00 ± 0.01	103.76 ± 11.69
spatial	0.27 ± 0.01	1.00 ± 0.00	79.60 ± 6.88

video 1, video 2, video 3.


Camera Environment

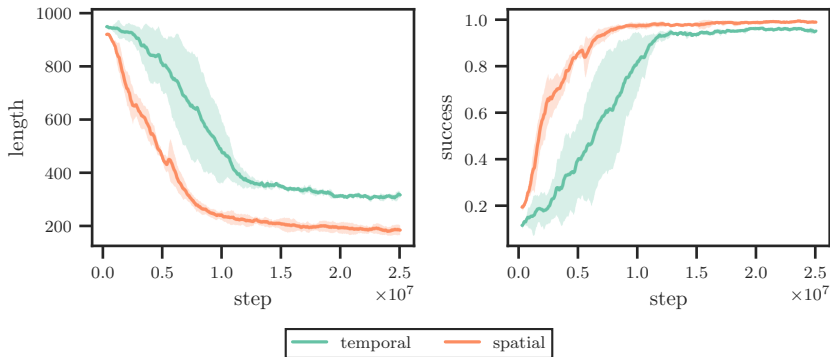
Agent	SPL	Success	Length
random	0.04 ± 0.00	0.62 ± 0.03	545.09 ± 56.25
greedy	0.12 ± 0.01	0.97 ± 0.01	255.60 ± 10.44
exhaustive	0.37 ± 0.00	1.00 ± 0.00	67.03 ± 0.00
human	0.68 ± 0.08	1.00 ± 0.00	38.10 ± 5.72
temporal	0.70 ± 0.02	1.00 ± 0.00	42.36 ± 2.05
spatial	0.66 ± 0.03	1.00 ± 0.00	42.90 ± 1.73


video 1, video 2, video 3.

Experiment II: Scaling to Larger Search Spaces

- ▶ Real-world search tasks usually have large search spaces.
- ▶ Stronger demands on memory:
 - ▶ Remember visited positions.
 - ▶ Remember appearance of environment.
- ▶ Compare memories on 10×10 , 15×15 , and 20×20 versions of gaussian environment.

10×10 

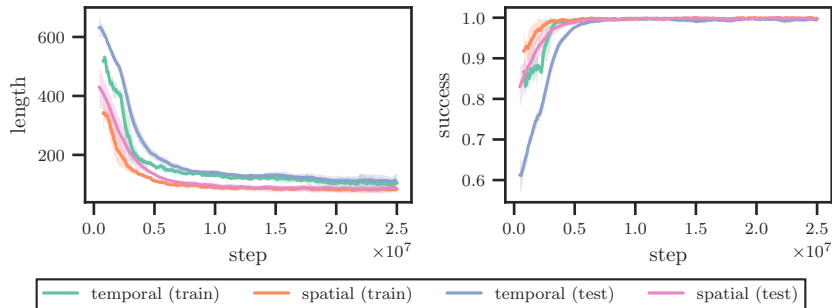
15×15 

20×20 


Experiment III: Generalization From Limited Samples

- ▶ Real-world tasks usually have limited training samples.
- ▶ Train on 500, 1 000, 5 000 and 10 000 samples of terrain environment.
- ▶ Test on held out samples from full distribution.

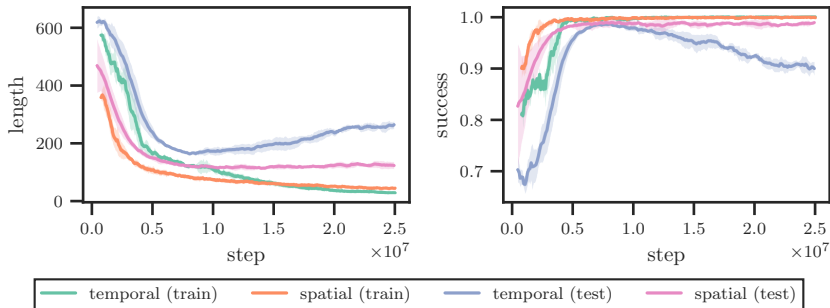
10000 samples




5000 samples



1000 samples



500 samples



Conclusion

- ▶ Architecture:
 - ▶ Spatial memory: architecture scales to larger search spaces and generalizes better.
 - ▶ Temporal memory: sufficient (and better) for smaller search spaces.
- ▶ Approach:
 - ▶ Search performance: better than simple baselines, comparable to human, worse than handcrafted.
 - ▶ Sample efficiency: relatively many samples needed even for simple environments.

Future Work

- ▶ Improvements to approach.
 - ▶ Neural network architecture.
 - ▶ Reinforcement learning algorithm.
 - ▶ Reward signal design.
- ▶ Evaluate on realistic search scenarios.
- ▶ Formal verification (for security-critical applications).

References I

- [1] A. Andreopoulos and J. K. Tsotsos, "A theory of active object localization," in *2009 IEEE 12th International Conference on Computer Vision*, pp. 903–910.
15 citations (Crossref) [2022-05-19] ISSN: 2380-7504.
- [2] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, "Human-level control through deep reinforcement learning," vol. 518, no. 7540, pp. 529–533.
Number: 7540 Publisher: Nature Publishing Group.
- [3] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. van den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, S. Dieleman, D. Grewe, J. Nham, N. Kalchbrenner, I. Sutskever, T. Lillicrap, M. Leach, K. Kavukcuoglu, T. Graepel, and D. Hassabis, "Mastering the game of go with deep neural networks and tree search," vol. 529, no. 7587, pp. 484–489.
Number: 7587 Publisher: Nature Publishing Group.

References II

- [4] O. Vinyals, I. Babuschkin, W. M. Czarnecki, M. Mathieu, A. Dudzik, J. Chung, D. H. Choi, R. Powell, T. Ewalds, P. Georgiev, J. Oh, D. Horgan, M. Kroiss, I. Danihelka, A. Huang, L. Sifre, T. Cai, J. P. Agapiou, M. Jaderberg, A. S. Vezhnevets, R. Leblond, T. Pohlen, V. Dalibard, D. Budden, Y. Sulsky, J. Molloy, T. L. Paine, C. Gulcehre, Z. Wang, T. Pfaff, Y. Wu, R. Ring, D. Yogatama, D. Wünsch, K. McKinney, O. Smith, T. Schaul, T. Lillicrap, K. Kavukcuoglu, D. Hassabis, C. Apps, and D. Silver, "Grandmaster level in StarCraft II using multi-agent reinforcement learning," vol. 575, no. 7782, pp. 350–354.
Number: 7782 Publisher: Nature Publishing Group.
- [5] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, "Planning and acting in partially observable stochastic domains," vol. 101, no. 1, pp. 99–134.
- [6] V. Mnih, N. Heess, A. Graves, and k. kavukcuoglu, "Recurrent models of visual attention," in *Advances in Neural Information Processing Systems*, vol. 27, Curran Associates, Inc.
- [7] P. Mirowski, R. Pascanu, F. Viola, H. Soyer, A. J. Ballard, A. Banino, M. Denil, R. Goroshin, L. Sifre, K. Kavukcuoglu, D. Kumaran, and R. Hadsell, "Learning to navigate in complex environments,"
- [8] Y. Zhu, R. Mottaghi, E. Kolve, J. J. Lim, A. Gupta, L. Fei-Fei, and A. Farhadi, "Target-driven visual navigation in indoor scenes using deep reinforcement learning," in *2017 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 3357–3364.
474 citations (Crossref) [2022-05-19].


References III

- [9] J. C. Caicedo and S. Lazebnik, “Active object localization with deep reinforcement learning,”
- [10] F.-C. Ghesu, B. Georgescu, Y. Zheng, S. Grbic, A. Maier, J. Hornegger, and D. Comaniciu, “Multi-scale deep reinforcement learning for real-time 3d-landmark detection in CT scans,” vol. 41, no. 1, pp. 176–189.
- [11] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,”
- [12] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, “Deep reinforcement learning that matters,” vol. 32, no. 1.
Number: 1.
- [13] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” vol. 9, no. 8, pp. 1735–1780.
Conference Name: Neural Computation.
- [14] M. Hausknecht and P. Stone, “Deep recurrent q-learning for partially observable MDPs,”
- [15] V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. P. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, “Asynchronous methods for deep reinforcement learning,”
- [16] E. Parisotto and R. Salakhutdinov, “Neural map: Structured memory for deep reinforcement learning,”
- [17] S. Gupta, V. Tolani, J. Davidson, S. Levine, R. Sukthankar, and J. Malik, “Cognitive mapping and planning for visual navigation,”


References IV

- [18] P. Anderson, A. Chang, D. S. Chaplot, A. Dosovitskiy, S. Gupta, V. Koltun, J. Kosecka, J. Malik, R. Mottaghi, M. Savva, and A. R. Zamir, “On evaluation of embodied navigation agents,”

Search Paths I




Environment sample




Random baseline

Search Paths II




Environment sample




Greedy baseline

Search Paths III




Environment sample




Exhaustive baseline

Search Paths IV




Environment sample




Handcrafted baseline

Search Paths V




Environment sample

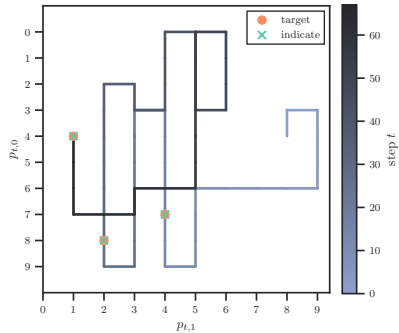


Temporal memory

Search Paths VI




Environment sample



Spatial memory

Memory Viualization



PCA decomposition of spatial memory after episode.